

# Lab 19

2022-10-26

## Trees

```
# Classification and Regression Trees (CART)
```

```
# Look at data!
```

```
head(titanic_train)
```

```
## PassengerId Survived Pclass
```

```
## 1 1 0 3
```

```
## 2 2 1 1
```

```
## 3 3 1 3
```

```
## 4 4 1 1
```

```
## 5 5 0 3
```

```
## 6 6 0 3
```

```
## Name Sex Age SibSp Parch
```

```
## 1 Braund, Mr. Owen Harris male 22 1 0
```

```
## 2 Cumings, Mrs. John Bradley (Florence Briggs Thayer) female 38 1 0
```

```
## 3 Heikkinen, Miss. Laina female 26 0 0
```

```
## 4 Futrelle, Mrs. Jacques Heath (Lily May Peel) female 35 1 0
```

```
## 5 Allen, Mr. William Henry male 35 0 0
```

```
## 6 Moran, Mr. James male NA 0 0
```

```
## Ticket Fare Cabin Embarked
```

```
## 1 A/5 21171 7.2500 S
```

```
## 2 PC 17599 71.2833 C85 C
```

```
## 3 STON/O2. 3101282 7.9250 S
```

```
## 4 113803 53.1000 C123 S
```

```
## 5 373450 8.0500 S
```

```
## 6 330877 8.4583 Q
```

```
?titanic_train
```

```
titanic_train$Survived = as.factor(titanic_train$Survived)
```

```
titanic_train %>%
```

```
ggpairs(columns = c("Pclass",
```

```
"Sex",
```

```
"Age",
```

```
"Fare"),
```

```
mapping = aes(color = Survived))
```

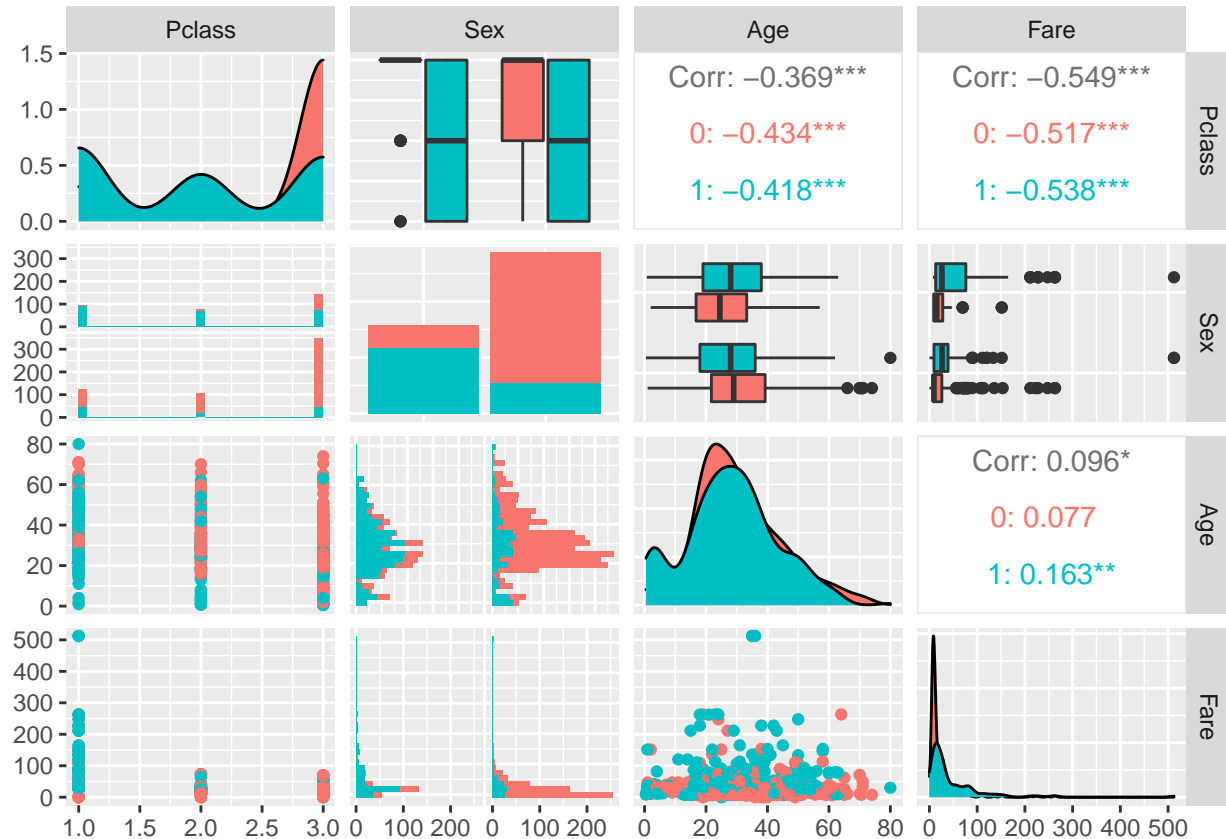
```
## Warning in ggally_statistic(data = data, mapping = mapping, na.rm = na.rm, :
```

```
## Removed 177 rows containing missing values
```


```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

```
## Warning: Removed 177 rows containing non-finite values (stat_boxplot).
```

```
## Warning: Removed 177 rows containing missing values (geom_point).
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
## Warning: Removed 177 rows containing non-finite values (stat_bin).
## Warning: Removed 177 rows containing non-finite values (stat_density).
## Warning in ggally_statistic(data = data, mapping = mapping, na.rm = na.rm, :
## Removed 177 rows containing missing values
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
## Warning: Removed 177 rows containing missing values (geom_point).
```



```
# Decision tree model
treeModel = rpart(Survived ~ Pclass + Age + Fare + Sex, data=titanic_train, method="class", xval=5, cp=
par(cex = 0.6)
plot(treeModel,margin=.05)
text(treeModel,use.n=TRUE)
```



```
## Call:
## rpart(formula = Survived ~ Pclass + Age + Fare + Sex, data = titanic_train,
##       method = "class", xval = 5, cp = 0.005)
##       n= 891
##
##              CP nsplit rel error      xerror      xstd
## 1 0.4444444444      0 1.0000000 1.0000000 0.04244576
## 2 0.030701754      1 0.5555556 0.5555556 0.03574957
## 3 0.023391813      3 0.4941520 0.5116959 0.03467453
## 4 0.011695906      4 0.4707602 0.5233918 0.03497048
## 5 0.006578947      7 0.4327485 0.5146199 0.03474917
## 6 0.005847953     13 0.3918129 0.4883041 0.03406141
## 7 0.005000000     15 0.3801170 0.4795322 0.03382394
##
## Variable importance
##      Sex      Fare Pclass      Age
##      49       26      18        8
##
## Node number 1: 891 observations,      complexity param=0.4444444
## predicted class=0 expected loss=0.3838384 P(node) =1
## class counts:      549      342
## probabilities: 0.616 0.384
## left son=2 (577 obs) right son=3 (314 obs)
## Primary splits:
##      Sex      splits as RL,      improve=124.426300, (0 missing)
##      Pclass < 2.5      to the right, improve= 43.781830, (0 missing)
##      Fare < 10.48125 to the left, improve= 37.941940, (0 missing)
##      Age < 6.5      to the right, improve= 8.814172, (177 missing)
## Surrogate splits:
##      Fare < 77.6229 to the left, agree=0.679, adj=0.089, (0 split)
##
## Node number 2: 577 observations,      complexity param=0.02339181
```

```

## predicted class=0 expected loss=0.1889081 P(node) =0.647587
## class counts: 468 109
## probabilities: 0.811 0.189
## left son=4 (553 obs) right son=5 (24 obs)
## Primary splits:
## Age < 6.5 to the right, improve=10.78893, (124 missing)
## Fare < 26.26875 to the left, improve=10.21672, (0 missing)
## Pclass < 1.5 to the right, improve=10.01914, (0 missing)
##
## Node number 3: 314 observations, complexity param=0.03070175
## predicted class=1 expected loss=0.2579618 P(node) =0.352413
## class counts: 81 233
## probabilities: 0.258 0.742
## left son=6 (144 obs) right son=7 (170 obs)
## Primary splits:
## Pclass < 2.5 to the right, improve=31.163130, (0 missing)
## Fare < 48.2 to the left, improve=10.114210, (0 missing)
## Age < 12 to the left, improve= 1.891684, (53 missing)
## Surrogate splits:
## Fare < 25.69795 to the left, agree=0.799, adj=0.563, (0 split)
## Age < 18.5 to the left, agree=0.564, adj=0.049, (0 split)
##
## Node number 4: 553 observations, complexity param=0.006578947
## predicted class=0 expected loss=0.1681736 P(node) =0.620651
## class counts: 460 93
## probabilities: 0.832 0.168
## left son=8 (433 obs) right son=9 (120 obs)
## Primary splits:
## Pclass < 1.5 to the right, improve=11.083720, (0 missing)
## Fare < 26.26875 to the left, improve=10.532060, (0 missing)
## Age < 24.75 to the left, improve= 1.235487, (124 missing)
## Surrogate splits:
## Fare < 26.26875 to the left, agree=0.911, adj=0.592, (0 split)
##
## Node number 5: 24 observations, complexity param=0.005847953
## predicted class=1 expected loss=0.3333333 P(node) =0.02693603
## class counts: 8 16
## probabilities: 0.333 0.667
## left son=10 (14 obs) right son=11 (10 obs)
## Primary splits:
## Pclass < 2.5 to the right, improve=3.8095240, (0 missing)
## Fare < 20.825 to the right, improve=2.6666670, (0 missing)
## Age < 1.5 to the right, improve=0.6095238, (0 missing)
## Surrogate splits:
## Age < 0.96 to the right, agree=0.708, adj=0.3, (0 split)
## Fare < 64.37915 to the left, agree=0.667, adj=0.2, (0 split)
##
## Node number 6: 144 observations, complexity param=0.03070175
## predicted class=0 expected loss=0.5 P(node) =0.1616162
## class counts: 72 72
## probabilities: 0.500 0.500
## left son=12 (27 obs) right son=13 (117 obs)
## Primary splits:
## Fare < 23.35 to the right, improve=10.051280, (0 missing)

```

```

##      Age < 38.5      to the right, improve= 3.875163, (42 missing)
##
## Node number 7: 170 observations
##   predicted class=1 expected loss=0.05294118 P(node) =0.1907969
##   class counts:      9   161
##   probabilities: 0.053 0.947
##
## Node number 8: 433 observations
##   predicted class=0 expected loss=0.1154734 P(node) =0.4859708
##   class counts:    383    50
##   probabilities: 0.885 0.115
##
## Node number 9: 120 observations,      complexity param=0.006578947
##   predicted class=0 expected loss=0.3583333 P(node) =0.1346801
##   class counts:      77    43
##   probabilities: 0.642 0.358
##   left son=18 (22 obs) right son=19 (98 obs)
##   Primary splits:
##     Age < 53          to the right, improve=3.464646, (21 missing)
##     Fare < 26.14375 to the left,  improve=2.801515, (0 missing)
##
## Node number 10: 14 observations
##   predicted class=0 expected loss=0.4285714 P(node) =0.01571268
##   class counts:      8     6
##   probabilities: 0.571 0.429
##
## Node number 11: 10 observations
##   predicted class=1 expected loss=0 P(node) =0.01122334
##   class counts:      0    10
##   probabilities: 0.000 1.000
##
## Node number 12: 27 observations
##   predicted class=0 expected loss=0.1111111 P(node) =0.03030303
##   class counts:     24     3
##   probabilities: 0.889 0.111
##
## Node number 13: 117 observations,      complexity param=0.01169591
##   predicted class=1 expected loss=0.4102564 P(node) =0.1313131
##   class counts:     48    69
##   probabilities: 0.410 0.590
##   left son=26 (91 obs) right son=27 (26 obs)
##   Primary splits:
##     Age < 16.5        to the right, improve=2.468587, (34 missing)
##     Fare < 7.8875     to the right, improve=2.032527, (0 missing)
##   Surrogate splits:
##     Fare < 20.8       to the left,  agree=0.747, adj=0.087, (34 split)
##
## Node number 18: 22 observations
##   predicted class=0 expected loss=0.1363636 P(node) =0.02469136
##   class counts:     19     3
##   probabilities: 0.864 0.136
##
## Node number 19: 98 observations,      complexity param=0.006578947
##   predicted class=0 expected loss=0.4081633 P(node) =0.1099888

```

```

##      class counts:      58      40
##      probabilities: 0.592 0.408
##      left son=38 (9 obs) right son=39 (89 obs)
##      Primary splits:
##          Fare < 26.14375 to the left, improve=3.301995, (0 missing)
##          Age < 36.5      to the right, improve=1.112992, (21 missing)
##
## Node number 26: 91 observations,      complexity param=0.01169591
##      predicted class=1 expected loss=0.4615385 P(node) =0.1021324
##      class counts:      42      49
##      probabilities: 0.462 0.538
##      left son=52 (54 obs) right son=53 (37 obs)
##      Primary splits:
##          Fare < 7.8875   to the right, improve=3.363902, (0 missing)
##          Age < 36.5      to the right, improve=1.661815, (31 missing)
##
## Node number 27: 26 observations
##      predicted class=1 expected loss=0.2307692 P(node) =0.0291807
##      class counts:      6      20
##      probabilities: 0.231 0.769
##
## Node number 38: 9 observations
##      predicted class=0 expected loss=0 P(node) =0.01010101
##      class counts:      9      0
##      probabilities: 1.000 0.000
##
## Node number 39: 89 observations,      complexity param=0.006578947
##      predicted class=0 expected loss=0.4494382 P(node) =0.09988777
##      class counts:      49      40
##      probabilities: 0.551 0.449
##      left son=78 (76 obs) right son=79 (13 obs)
##      Primary splits:
##          Fare < 27.1354 to the right, improve=4.7919070, (0 missing)
##          Age < 43      to the right, improve=0.8888889, (17 missing)
##
## Node number 52: 54 observations,      complexity param=0.01169591
##      predicted class=0 expected loss=0.4259259 P(node) =0.06060606
##      class counts:      31      23
##      probabilities: 0.574 0.426
##      left son=104 (33 obs) right son=105 (21 obs)
##      Primary splits:
##          Fare < 14.8729 to the left, improve=2.563252, (0 missing)
##          Age < 23.5     to the left, improve=1.314848, (11 missing)
##
## Node number 53: 37 observations
##      predicted class=1 expected loss=0.2972973 P(node) =0.04152637
##      class counts:      11      26
##      probabilities: 0.297 0.703
##
## Node number 78: 76 observations,      complexity param=0.006578947
##      predicted class=0 expected loss=0.3815789 P(node) =0.08529742
##      class counts:      47      29
##      probabilities: 0.618 0.382
##      left son=156 (7 obs) right son=157 (69 obs)

```


```

## Primary splits:
##   Fare < 29.85   to the left,  improve=0.8787730, (0 missing)
##   Age  < 43      to the right, improve=0.7179528, (15 missing)
##
## Node number 79: 13 observations
##   predicted class=1  expected loss=0.1538462  P(node) =0.01459035
##   class counts:      2    11
##   probabilities: 0.154 0.846
##
## Node number 104: 33 observations
##   predicted class=0  expected loss=0.3030303  P(node) =0.03703704
##   class counts:      23    10
##   probabilities: 0.697 0.303
##
## Node number 105: 21 observations,   complexity param=0.005847953
##   predicted class=1  expected loss=0.3809524  P(node) =0.02356902
##   class counts:      8    13
##   probabilities: 0.381 0.619
##   left son=210 (8 obs) right son=211 (13 obs)
##   Primary splits:
##     Fare < 17.6      to the right, improve=1.53937700, (0 missing)
##     Age  < 30         to the left,  improve=0.03809524, (6 missing)
##
## Node number 156: 7 observations
##   predicted class=0  expected loss=0.1428571  P(node) =0.007856341
##   class counts:      6    1
##   probabilities: 0.857 0.143
##
## Node number 157: 69 observations,   complexity param=0.006578947
##   predicted class=0  expected loss=0.4057971  P(node) =0.07744108
##   class counts:      41    28
##   probabilities: 0.594 0.406
##   left son=314 (62 obs) right son=315 (7 obs)
##   Primary splits:
##     Fare < 30.5979   to the right, improve=3.1739800, (0 missing)
##     Age  < 43         to the right, improve=0.8596491, (12 missing)
##
## Node number 210: 8 observations
##   predicted class=0  expected loss=0.375  P(node) =0.008978676
##   class counts:      5    3
##   probabilities: 0.625 0.375
##
## Node number 211: 13 observations
##   predicted class=1  expected loss=0.2307692  P(node) =0.01459035
##   class counts:      3    10
##   probabilities: 0.231 0.769
##
## Node number 314: 62 observations
##   predicted class=0  expected loss=0.3548387  P(node) =0.06958474
##   class counts:      40    22
##   probabilities: 0.645 0.355
##
## Node number 315: 7 observations
##   predicted class=1  expected loss=0.1428571  P(node) =0.007856341

```

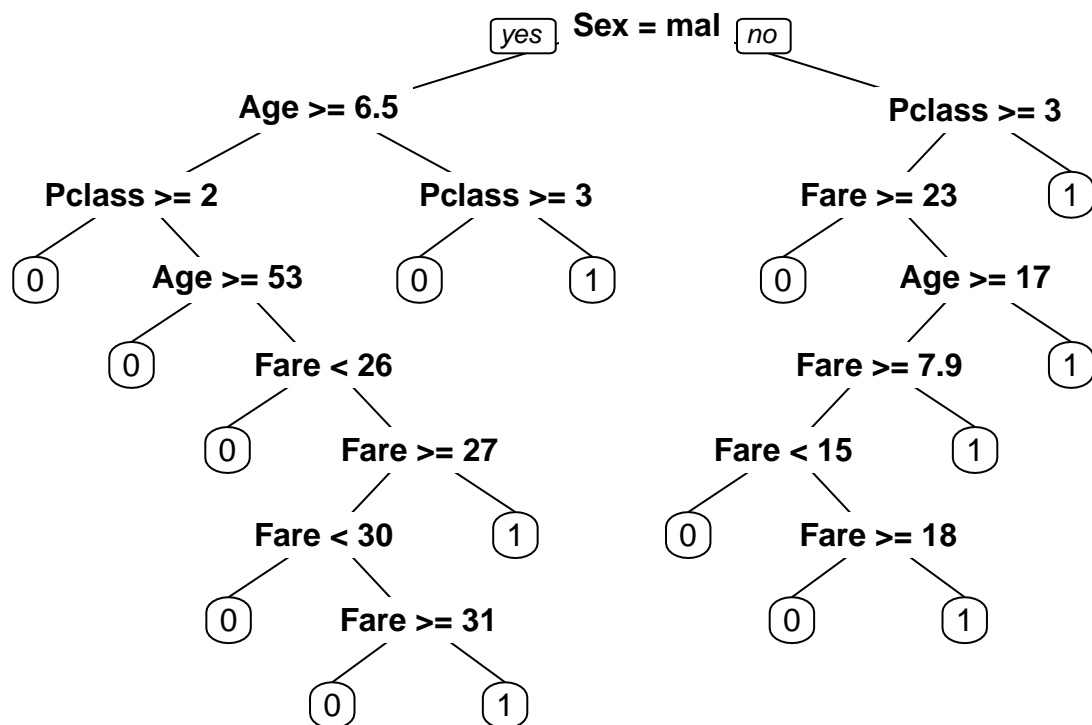
```
##      class counts:      1      6
##      probabilities: 0.143 0.857
```

```
par(cex = 1)
plotcp(treeModel)
```

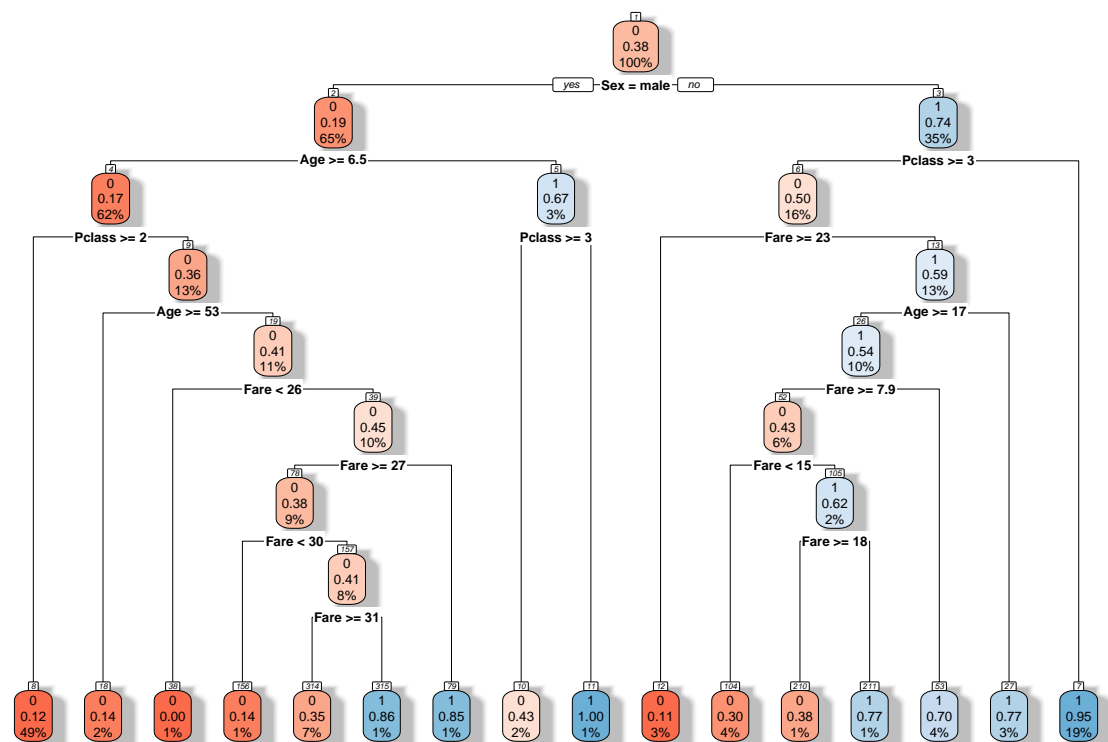


```
# Some nicer plots
prp(treeModel)
```





```
rpart.plot(treeModel, box.palette="RdBu", shadow.col="gray", nn=TRUE)
```




## Plotting decision boundaries

```
# Plotting decision boundaries
```

```
# Create a model with 2 predictors
treeModel = rpart(Survived ~ Pclass + Age, data=titanic_train, method="class", xval=5, cp=.00005)

# Plot decision boundaries
titanic_train %>%
  ggplot(aes(x=Pclass, y=Age)) +
  geom_jitter(aes(col=Survived), alpha=0.7) +
  geom_parttree(data = treeModel, aes(fill=Survived), alpha = 0.1) +
  theme_minimal()
```


## Warning: Removed 177 rows containing missing values (geom\_point).



```
# What happens if we tweak the complexity parameter a.k.a. cp
# cp = 0.01
treeModel = rpart(Survived ~ Pclass + Age, data=titanic_train, method="class", xval=5, cp=.01)

# Plot decision boundaries
titanic_train %>%
  ggplot(aes(x=Pclass, y=Age)) +
  geom_jitter(aes(col=Survived), alpha=0.7) +
  geom_parttree(data = treeModel, aes(fill=Survived), alpha = 0.1) +
  theme_minimal()
```


## Warning: Removed 177 rows containing missing values (geom\_point).



```
# cp = 0.1
treeModel = rpart(Survived ~ Pclass + Age, data=titanic_train, method="class", xval=5, cp=.1)

# Plot decision boundaries
titanic_train %>%
  ggplot(aes(x=Pclass, y=Age)) +
  geom_jitter(aes(col=Survived), alpha=0.7) +
  geom_parttree(data = treeModel, aes(fill=Survived), alpha = 0.1) +
  theme_minimal()

## Warning: Removed 177 rows containing missing values (geom_point).
```




```
# Let us split into training and test sets
## 75% of the sample size
train.index <- createDataPartition(titanic_train$Survived, p = .75, list = FALSE)
train <- titanic_train[ train.index,]
test  <- titanic_train[-train.index,]

# Train a model and check performance on the test set
treeModel = rpart(Survived ~ Pclass + Age, data=train, method="class", xval=5, cp=.01)

# Plot decision boundaries
train %>%
  ggplot(aes(x=Pclass, y=Age)) +
  geom_jitter(aes(col=Survived), alpha=0.7) +
  geom_parttree(data = treeModel, aes(fill=Survived), alpha = 0.1) +
  theme_minimal()

## Warning: Removed 125 rows containing missing values (geom_point).
```



```
surv_pred <- predict(treeModel, newdata=test, type='class')
surv_pred
```

```
##      5      6     13     14     17     19     20     24     28     29     33     36     39     42     43     46     48     50     51     56
##      0      0      0      0      0      0      0      1      1      0      0      1      0      0      0      0      0      0      0      1
##    57    59    60    66    73    77    84    86    91    94    97    99   101   108   111   128   136   140   142   153
##      0      1      0      0      0      0      1      0      0      0      0      0      0      0      0      0      0      1      0      0
##   154   158   162   167   173   176   179   180   181   184   187   188   195   196   200   210   218   219   221   232
##      0      0      0      1      0      0      0      0      0      1      0      0      0      1      0      1      0      1      0      0
##   234   236   237   239   242   243   244   247   257   263   264   265   271   278   282   283   284   306   307   312
##      0      0      0      0      0      0      0      0      1      1      1      0      1      0      0      0      0      1      1      1
##   314   317   324   327   335   336   345   351   353   355   357   361   365   366   370   377   380   383   386   393
##      0      0      0      0      1      0      0      0      0      0      1      0      0      0      1      0      0      0      0      0
##   394   400   403   408   413   419   420   421   426   427   430   433   438   457   459   462   469   475   476   482
##      1      0      0      1      1      0      0      0      0      0      0      0      0      0      0      0      0      0      1      0
##   487   488   503   512   517   518   520   523   524   526   535   536   540   543   549   551   558   562   563   574
##      1      1      0      0      0      0      0      0      0      0      0      1      1      0      0      1      1      0      0      0
##   577   580   584   585   586   588   593   598   599   602   604   610   615   617   618   620   623   624   627   629
##      0      0      1      0      1      1      0      0      0      0      0      1      0      0      0      0      0      0      0      0
##   631   632   634   637   644   645   650   653   661   666   668   681   682   686   688   690   693   696   699   700
##      0      0      1      0      0      0      0      0      1      0      0      0      1      0      0      1      0      0      1      0
##   708   713   719   723   736   738   739   742   745   750   752   756   771   774   775   780   790   792   793   794
##      1      1      0      0      0      1      0      1      0      0      0      1      0      0      0      0      0      1      0      1
##   802   804   810   814   817   825   829   830   833   834   835   837   840   843   844   848   857   868   871   881
##      0      0      1      0      0      0      0      0      0      0      0      0      1      1      0      0      0      1      0      0
##   887   891
##      0      0
```

```
## Levels: 0 1
```

```
confusionMatrix(surv_pred, test$Survived)
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction  0    1
```

```
##           0 120  50
```

```
##           1  17  35
```

```
##
```

```
##           Accuracy : 0.6982
```

```
##           95% CI : (0.6332, 0.7578)
```

```
## No Information Rate : 0.6171
```

```
## P-Value [Acc > NIR] : 0.007179
```

```
##
```

```
##           Kappa : 0.3106
```

```
##
```

```
## McNemar's Test P-Value : 9.252e-05
```

```
##
```

```
##           Sensitivity : 0.8759
```

```
##           Specificity : 0.4118
```

```
## Pos Pred Value : 0.7059
```

```
## Neg Pred Value : 0.6731
```

```
## Prevalence : 0.6171
```

```
## Detection Rate : 0.5405
```

```
## Detection Prevalence : 0.7658
```

```
## Balanced Accuracy : 0.6438
```

```
##
```

```
## 'Positive' Class : 0
```

```
##
```

```
# plotting the confusion matrix
```

```
trueAndPredFr <- data.frame(surv_pred, test$Survived)
```

```
confMat <- conf_mat(trueAndPredFr, truth=test.Survived, estimate=surv_pred)
```


```
autoplot(confMat, type = "heatmap") +
```

```
  scale_fill_gradient(low="#D6EAF8",high = "#2E86C1") +
```

```
  theme(legend.position = "right")
```

```
## Scale for 'fill' is already present. Adding another scale for 'fill', which
```

```
## will replace the existing scale.
```



*# Take home exercise: Can you tune `cp` using a train and validation set, and then  
# test the performance on a test set?*